Indexing

Agenda

Why & what of indexes?
How indexes work
Indexes on multiple columns
Cons of Indexes
Indexes on Strings.

select * from film where film_id = 100;

↓

TABLE SCAN : row by row ⟶ O(N)

B1 B2 B3 B4

↑

202

HM : film_id, address.

↓

O(1).

Index , indexing

select * from users where name = "MOHIT";

create an index [key : name , value : row.]

We can have duplicate name

HM < String, int > ✗

HM < String, List >

Mohit	3, 7, 13, 18
Akash	1, 10, 16, 18

Select * from film where release-year between (2016, 2023);

[Key: release-year, value: List< >]

2016
2017
2018
2019
2020
⋮

2019
2023

range

HM's are helpful → ordered

$log_2 N$

HM + ordered

TreeMap / ordered-map

BBST: Balanced Binary Search Tree.

O(H)

$O(log_2 N)$

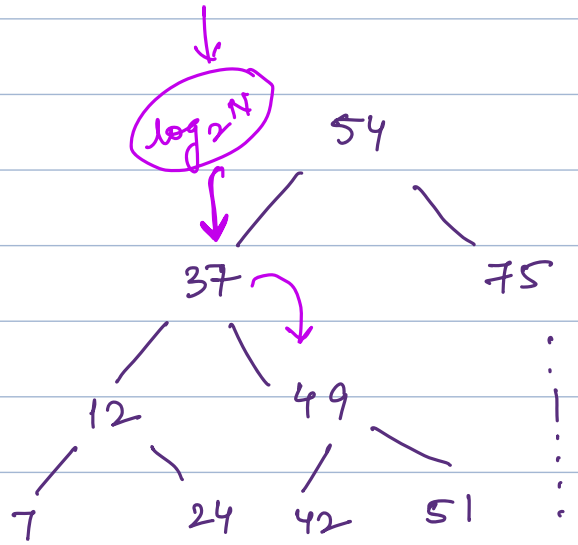

inorder

B / B+ Trees.

decrease height

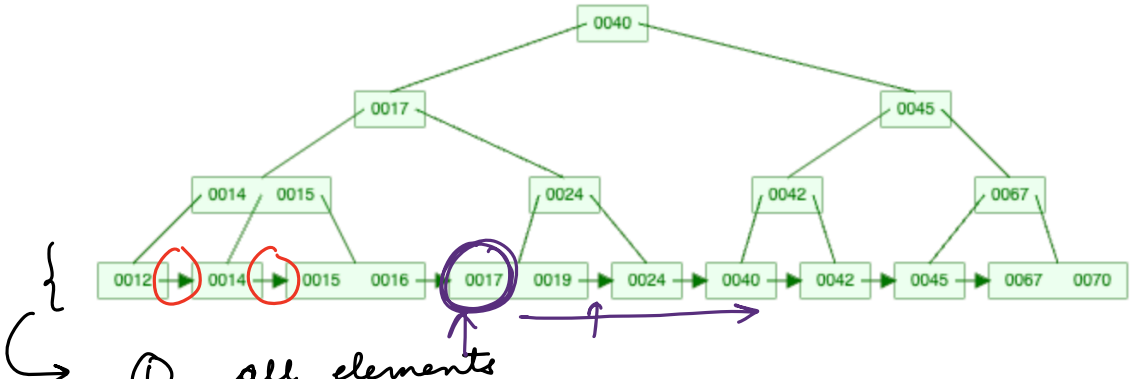

① all elements

② leaf node maintains a pointer to next leaf Node.

mysql : tables : indexes : clustered indexes

↑

indexes on primary keys

id	a	b	c
1	Ag	B	C
2	Au	Be	Co
3	Al	Br	Cr
4	Ar	Br	Cd
5	Ar	Br	C
6	Ag	B	Co
7	At	Bi	Ce
8	Al	B	C
9	Al	B	Cd
10	Ar	B	Cd

Ag Au Al Ar

id	a	b	c
1	Ag	B	C
2	Au	Be	Co
3	Al	Br	Cr
4	Ar	Br	Cd

primary key

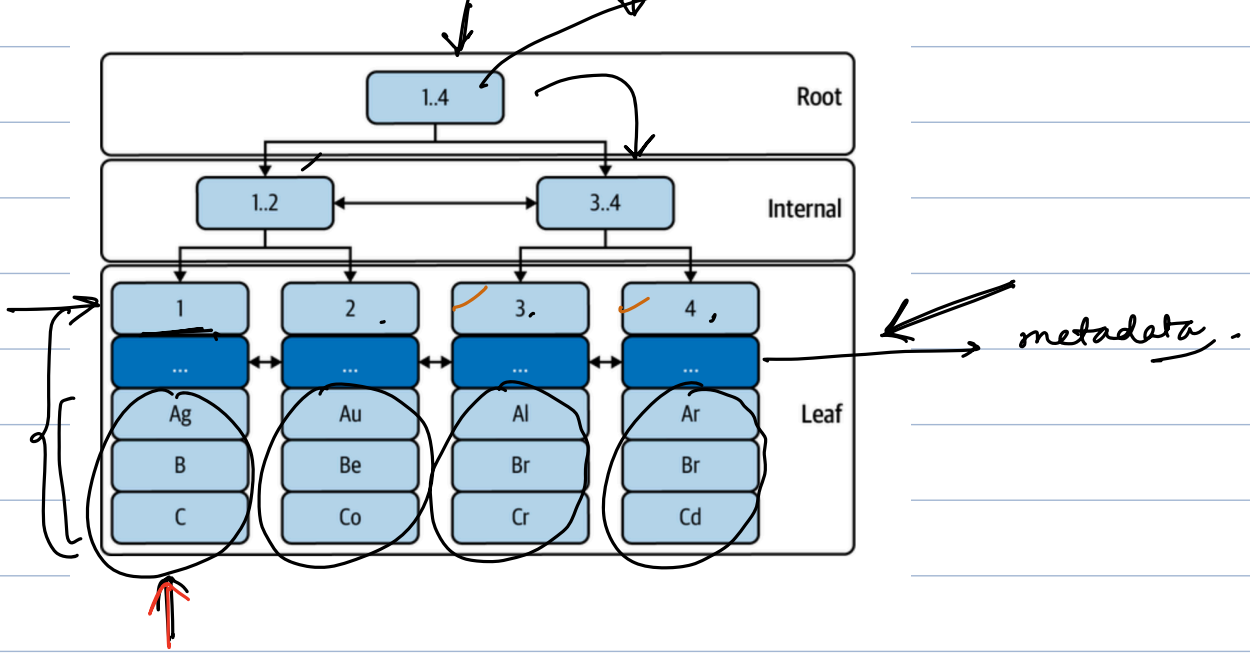

select * from film
order by film_id;

We can create our own indexes: Secondary, Non-clustered index.

select * from Symbol where a = 'Ag' and b = 'Be';

column → PK .

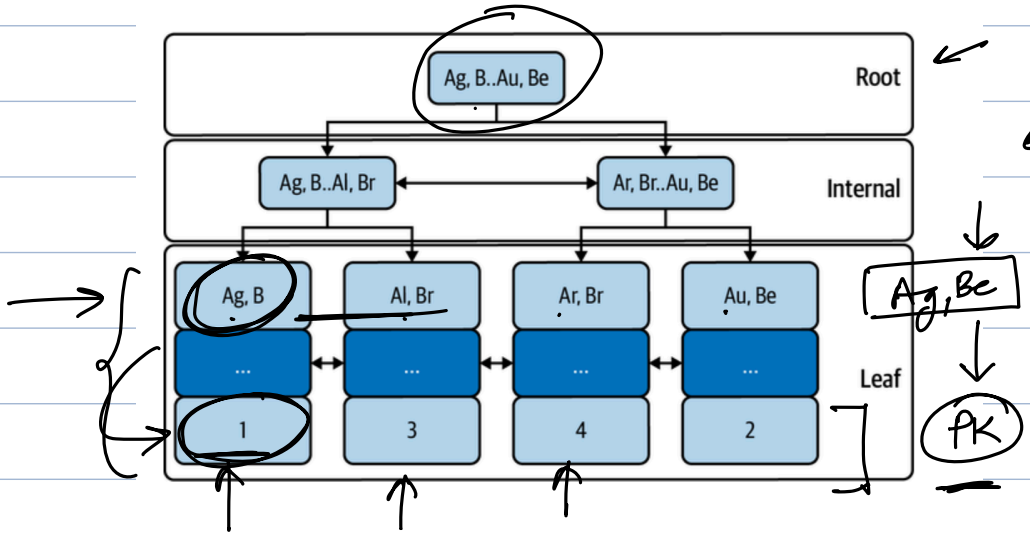

Ag, Al, Ar, Au

Multiple columns

where film_id = ___ ;

order by film_id ;

customers

email, PAN

where email = "___"

& PAN = ___ ;

indexes on PAN → unique

complete table scan.

brand

category

brand, category .

brand, category .X

rows = 10,000

brand = 10 , categories = 100

10000 / 10 | 1000

10000 / 100 | 100 prodn.

(A, B)

(B, A)

] diff indexes

order matters -

{ marks, id
id, marks }

idx ⟶ (f_name, last_name)

idx ⟶ (actor_id) ✗
PRIMARY

complete data

Select * from actors
where f_name = ' '
& l_name = ' ',

first_name

Select ^ from actors
where f_name = ' '
& l_name = ' '

↓
covering indexes .

I actor_id
II actor_id, email
III last_name, email

LEFT PREFIX

actor_id ⟶ I
last_name ⟶ III
last_name, actor_id ⟶ (I)
email ⟶ None of the above .

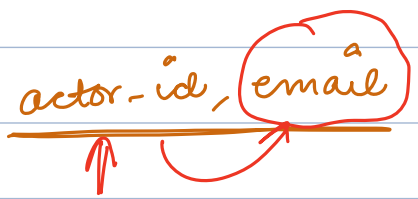

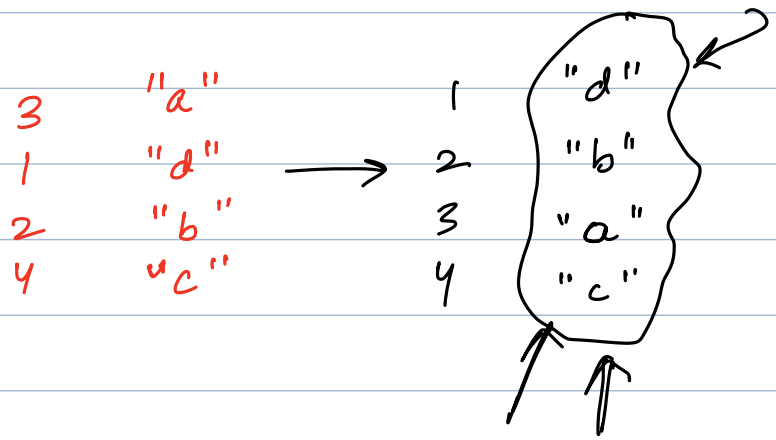

a, b, c, d, e

I a

II a, c

III a, c, b

IV a, b, c

c and b X

c X

a and b

first_N, L_Name

first_Name

X

useless index

first_name

last_name

new idx will be required.

where a = ____ and b = ____ ;

a = ____ OR b = ____

a

b → x True

(a, b)

(a, c)

a or b

CONS OF INDEXING

① indexes occupy space

② C R U D → writes can become slower
↓
update

a — d

aman bro Cat Daman
↓
Mount

12 columns
↓
12 indexes ✓

insert

[BIT MAP INDEX]

low cardinality :- no of unique values.
↓

selectivity : no of unique / total → 0-1

drink
{ coffee
  tea

1 M row
10^6 rows.

2

2 / 1 M

→ 0.5 M → O(n/2)
O(n)

select * from
bevy where drink = "tea";

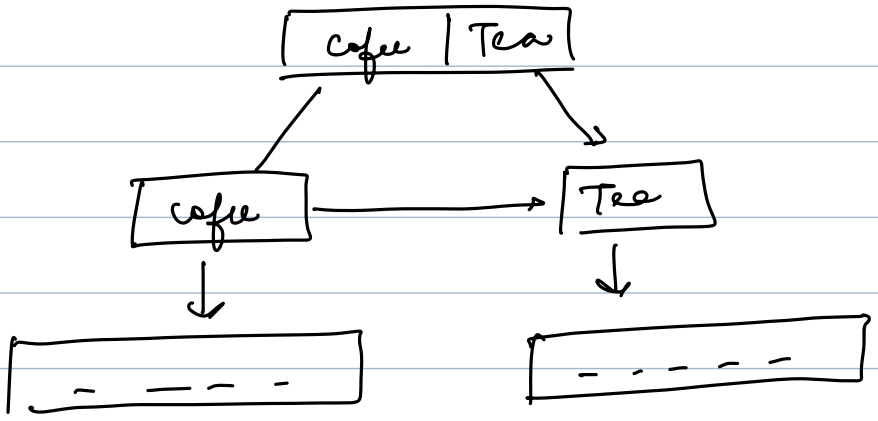
Coffee | Tea
Coffee
Tea

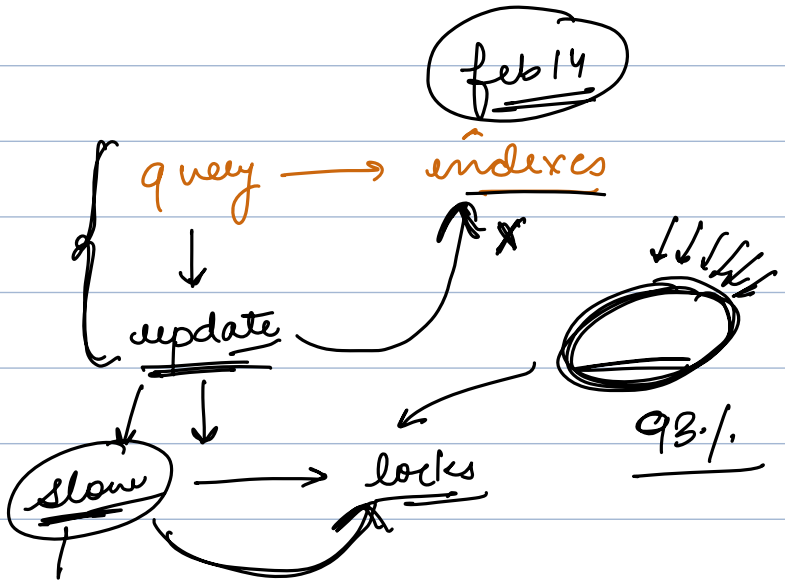
feb 14
query
indexes
update
slow
locks
93%

INDEXING ON STRINGS

① bad in comparison

② space

10 @gmail.com

where email = "   ";

an index ✓

10 M records

↳ 9M @gmail.com

create an idx_email on customer(email(7));

2B records

26 → 26 * 26 = 676

$26^4 = 4*10^5$

$26^5 = 10^7$

$26^6 = 3*10^8$

17576

$26^7 = 8*10^9$

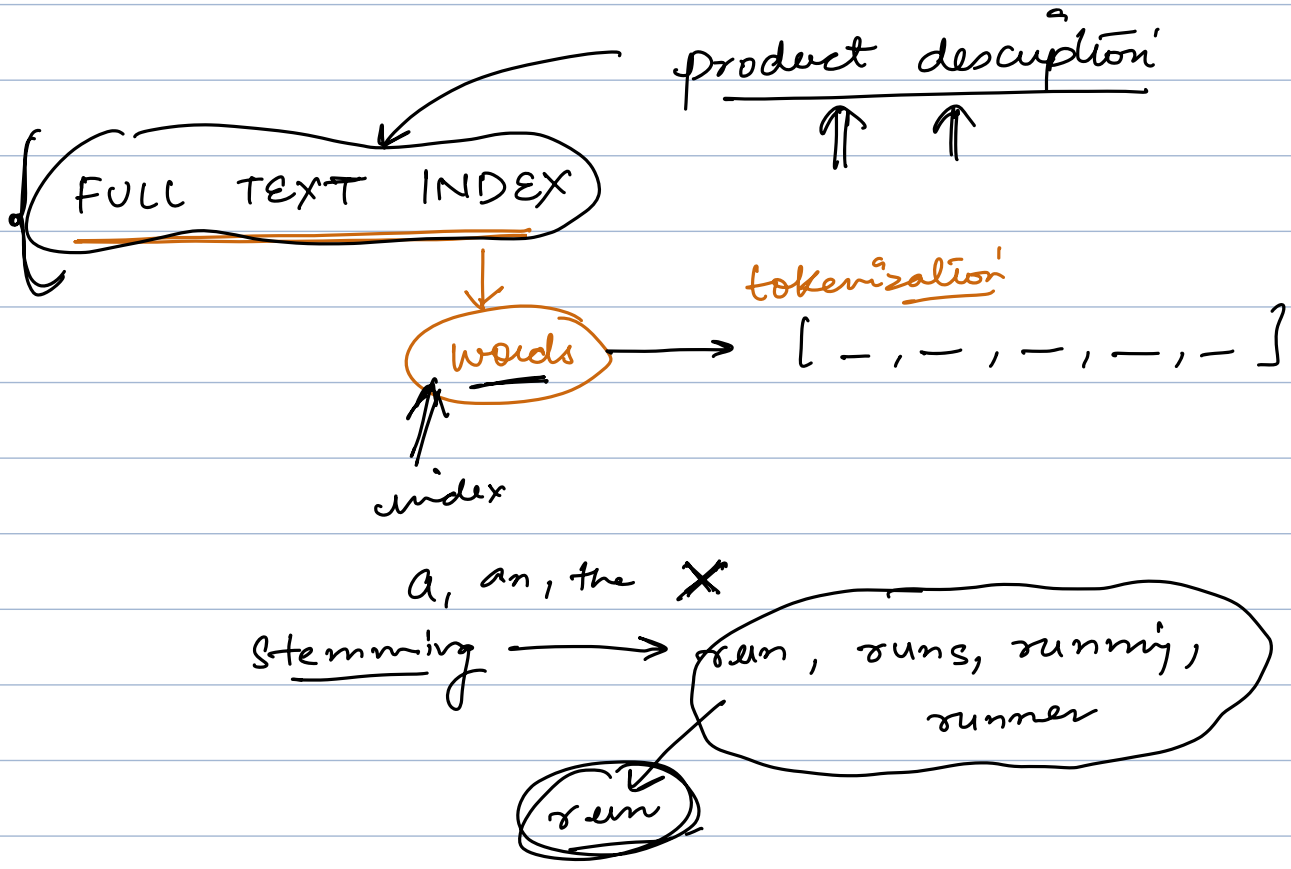
product description
FULL TEXT INDEX
tokenization
words
[_ , _ , _ , _ , _]
index
a, an, the
Stemming
run, runs, running, runner
run
MONITOR PERFORMANCE